

TRAFFIC SIGN RECOGNITION

Author: Jumana Mundichipparakkal

DATA SET SUMMARY & EXPLORATION

- 1- **Provide a basic summary of the data set. In the code, the analysis should be done using python, numpy and/or pandas methods rather than hardcoding results manually.**

I used the numpy library to calculate summary statistics of the traffic signs data set:

- The size of training set is : 34799
- The size of the validation set is : 4410
- The size of test set is : 12630
- The shape of a traffic sign image is : (32, 32, 3)
- The number of unique classes/labels in the data set is : 43

- 2- **Include an exploratory visualization of the dataset.**

Class Visualization:

Initially, I created a pandas dataframe which parsed the signal class data and shows the mapping between the ClassId and SignName. In order to visualize the unique signals, I identified unique indexes per class to show one signal image per class.

Data Distribution:

There are three sets of data provided: training, validation and test data.

I decided to plot histograms/distribution of image count per class in all the three sets. The plot clearly shows the huge variation in the image count per class, which could be an issue for obtaining a fair classification process done.

From Figure1: Training set has the lowest of 180 and highest of ~2000

From Figure2: Validation set has the lowest of 180 and highest of ~250

From Figure3: Test set has the lowest of 180 and highest of ~720

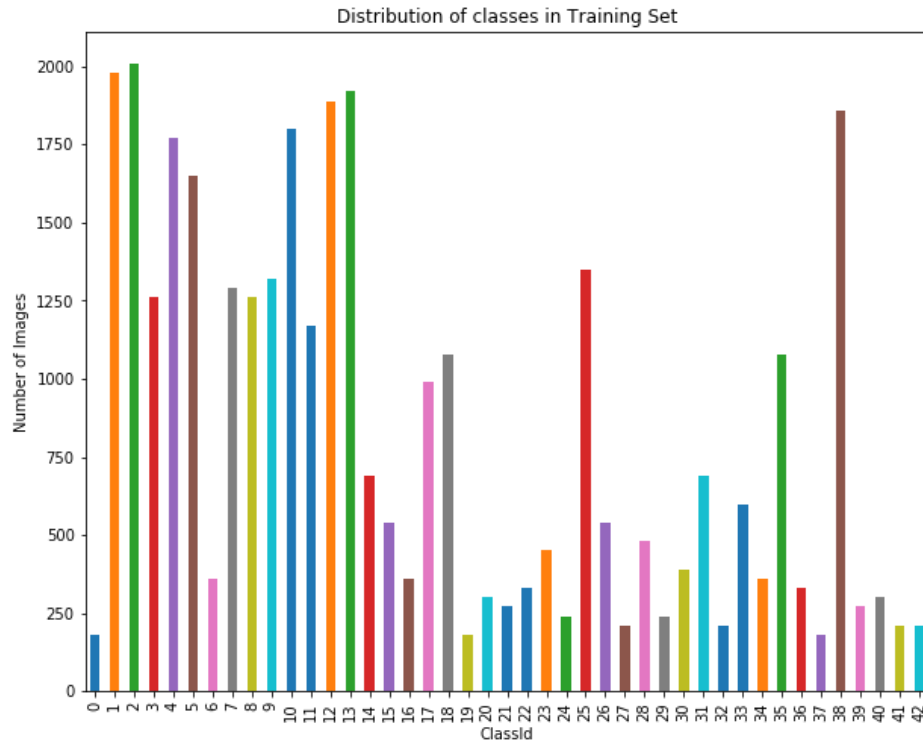


Figure 1 Training Set Data Distribution

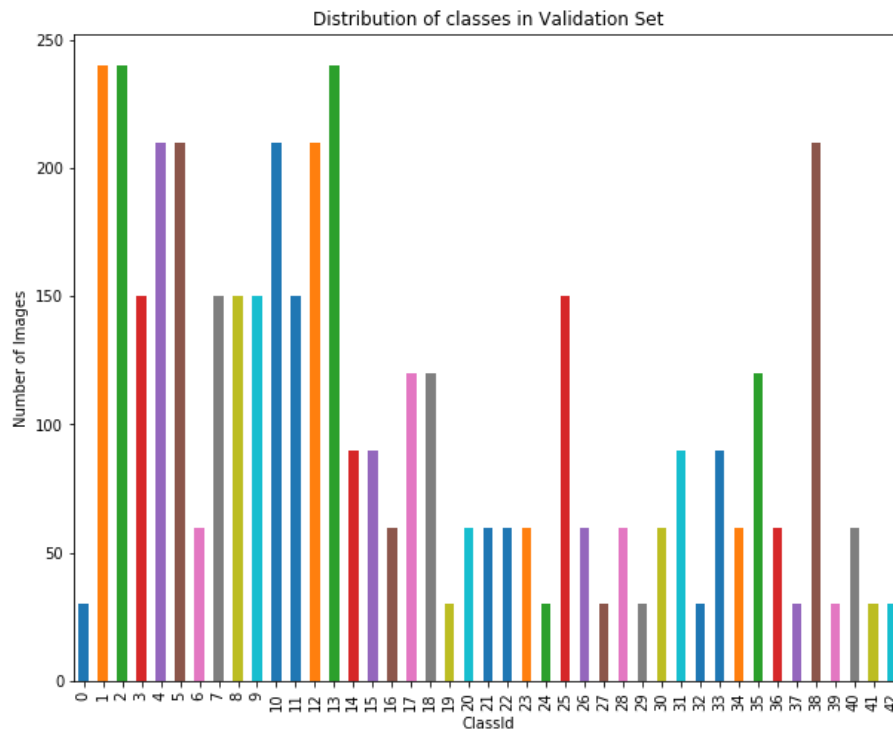


Figure 2 Validation Set Data Distribution

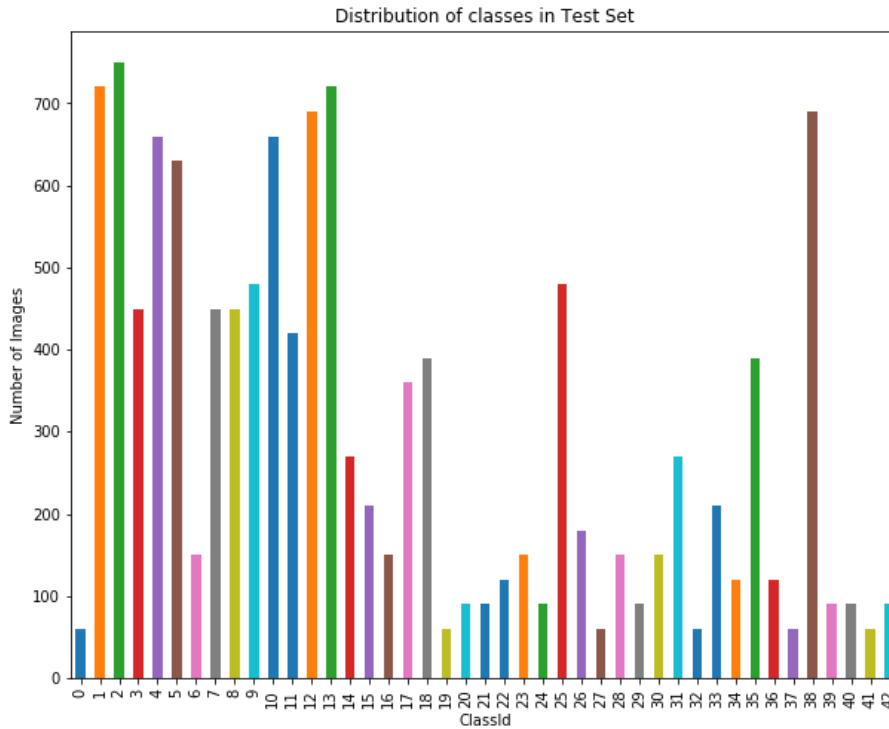


Figure 3 Test Set Data Distribution

DESIGN AND TEST A MODEL ARCHITECTURE

1. Describe how you preprocessed the image data. preprocessing technique.

For preprocessing, I initially tried to normalize the RGB color space by histogram equalization and adaptive normalization (functions that I used for this are still in the notebook). Later I realized that color should not matter a lot in identifying traffic signs and reducing the color channels will help in training the data set quicker without a GPU as my GPU crashed during the work. I decided to go for RGB-YUV color channel conversion, and then use the Y channel only, for my final solution.

Speed limit (20km/h)



Figure 4 Some of the Image Transformation Steps Experimented with

Finally Adopted Preprocessing Technique:

My final dataset pre-processing consisted of:

1. Converting to Y-channel
2. Normalizing the data to the range (-0.5, 0.5) - This was done using the line of code $X_train_y_norm = -0.5 + (X_train_y - 0)/(255-0)$. Normalized data can make the training faster and reduce the chance of getting stuck in local optima
3. Augmentation of Images using Image Transformation Techniques:

Fundamental requirement for a deep learning process is having a lot of data to learn from. So the more the data available, the learning will be more accurate. Moreover, from the data distribution visualization, it was clear that images vary a lot between classes. This can cause a bias between images that are predicted wrong or identical to choose the class that has more representative images. To prevent such bias cases, additional images were generated per class to equalize the number in the training set. Additional images can be added by applying some image transformations on existing images to produce new images. A few of the techniques tried are shown in Figure 4. Initially, I tried adding mean number of images per class. By trial and error, I decided to go for 3000 images per class. For generating new images, I used rotation of existing images randomly choosing angles to rotate among -10, 10, -15 and 15 degrees.

I wanted to use other techniques like cropping images, scaling images etc but I did not have time to test the effects of all such affine transformations as my GPU crashed during the process. Code for these transformations I initially used with RGB image are provided. Due to time limitations, I decided to settle for only rotation of images for my final submitted solutions.




Figure 5 Augmentation to fill 3000 images per class

2. Describe what your final model architecture looks like including model type, layers, layer sizes, connectivity, etc.) Consider including a diagram and/or table describing the final model.

I use a convolutional neuronal network to classify the traffic signs. The input of the network is a 32x32x1 image (Y-channel image) and the output is the probability of each of the 43 possible traffic signs.

My final model consisted of the following layers:

Layer	Description	Input	Output
Convolution 5x5	1x1 stride, valid padding, RELU activation	32x32x1	28x28x32
Max pooling 2x2	2x2 stride, 2x2 window	28x28x32	14x14x32
Convolution 5x5	1x1 stride, valid padding, RELU activation	14x14x32	10x10x64
Max pooling 2x2	2x2 stride, 2x2 window	10x10x64	5x5x64
Convolution 3x3	1x1 stride, valid padding, RELU activation	5x5x64	3x3x128
Max pooling 2x2	2x2 stride, 2x2 window	3x3x128	2x2x128
Flatten	3 dimensions to 1 dimension	2x2x128	512
Fully Connected 512	Connect every neuron from layer above	512	120
Fully Connected 120	Connect every neuron from layer above	120	84
Fully Connected 84	Output is equal to the number of traffic signs in data set	84	43

3. Describe how you trained your model. The discussion can include the type of optimizer, the batch size, number of epochs and any hyperparameters such as learning rate.

Final training parameters are:

- EPOCHS = 10
- BATCH_SIZE = 128
- SIGMA = 0.1
- OPIMIZER: AdamOptimizer (learning rate = 0.001)

Results after training the model:

- Validation Accuracy = **96.2%**
- Test Accuracy = **94%**

4. Describe the approach taken for finding a solution and getting the validation set accuracy to be at least 0.93. Include in the discussion the results on the training, validation and test sets and where in the code these were calculated. Your approach may have been an iterative process, in which case, outline the steps you took to get to the final solution and why you chose those steps. Perhaps your solution involved an already well known implementation or architecture. In this case, discuss why you think the architecture is suitable for the current problem.

My first implementation was LeNet-5 shown in the udacity classroom. I modified it to work with the input shape of 32x32x3 for RGB images. It was a good starting point and I got a test accuracy of close to 90% only with that approach. Some of the logs of experiments done with Lenet-5 on RGB/YUV 3 channel images images with some preprocessing and transformations are as below.

Architecture	Description	Validation Accuracy
Lenet-5	- Histogram equalized and normalized RGB image - Augment images to mean count with only rotation	89.2
Lenet-5	- Histogram equalized and normalized RGB image - Augment images to 2000 with only rotation	86.4
Lenet-5	- Histogram equalized and normalized RGB image -Augment images to 2000 with rotation and cropping	86.9
Lenet-5	- Histogram equalized and normalized YUV image -Augment images to 2000 with rotation and cropping	90.8
Lenet-5	- Histogram equalized and normalized YUV image -Augment images to mean count with rotation and cropping	86.7
Lenet-5	- Histogram equalized and normalized YUV image -Augment images to 3000 with rotation and cropping	89.2

Later I modified the network for YUV image with only one channel. Finding that accuracy keeps varying for validation randomly and accuracy was high for validation, it looked like that data was overfitting. Therefore, I added an extra convolution layer as well as increased filter depths as shown in my final architecture. This time I improved my test accuracy close to 95%. Still seeing not a steady increase in validation accuracy after each epoch, I added dropout in the fully connected layers to prevent overfitting. I used dropout probability as 0.5.

Architecture	Description	Validation Accuracy
My_Lenet without dropout	- Y-channel Normalized Image to -0.5 to 0.5 - Augment images to 3000 with only rotation	95.7
My_Lenet With dropout at first two FCL	- Y-channel Normalized Image -0.5 to 0.5 - Augment images to 3000 with only rotation	96.2

Experimenting with Epochs:

Training for more than 10 epochs do not increase the validation accuracy. I trained the network for 30 and more epochs, but I get a slightly decreasing accuracy. So, I decided to keep training for epochs which saves training time on a CPU as well.

Google Lenet suggests improving further by using advanced techniques like inception. I got an idea about how such an architecture could improve but due to time constraints I did not go about experimenting on further modifying my architecture.

TEST A MODEL ON NEW IMAGES

1. Choose five German traffic signs found on the web and provide them in the report. For each image, discuss what quality or qualities might be difficult to classify.



Figure 6 Test Images chosen

I collected five samples of traffic signals as above in Figure 6.

The signs "speed limit 60km/h", "yield, stop", and "keep left" should be easy to detect, because they are clearly visible and there are lots of examples in the training set. The "road_work" sign should be a little bit tricky, because there are only parts of the sign visible.

2. Discuss the model's predictions on these new traffic signs and compare the results to predicting on the test set. At a minimum, discuss what the predictions were, the accuracy on these new predictions, and compare the accuracy to the accuracy on the test set (OPTIONAL: Discuss the results in more detail as described in the "Stand Out Suggestions" part of the rubric).



Figure 7 Predictions on Test Images

Image Input	Prediction
Speed Limit 30km/h	Speed Limit 30km/h
Left Turn	Left Turn
Road Work	Speed Limit 20km/h
Stop Sign	Stop
Yield Sign	Yield

4 of 5 correct = **80.0 %**

Predictions are shown in Figure 7. 4 of 5 were predicted right giving an accuracy of 80%, which is much lower than the accuracy of the test set (94 %). I guess that with low number of test images, it is hard to predict it all right.

3. Describe how certain the model is when predicting on each of the five new images by looking at the softmax probabilities for each prediction. Provide the top 5 softmax probabilities for each image along with the sign type of each probability. (OPTIONAL: as

described in the "Stand Out Suggestions" part of the rubric, visualizations can also be provided such as bar charts)

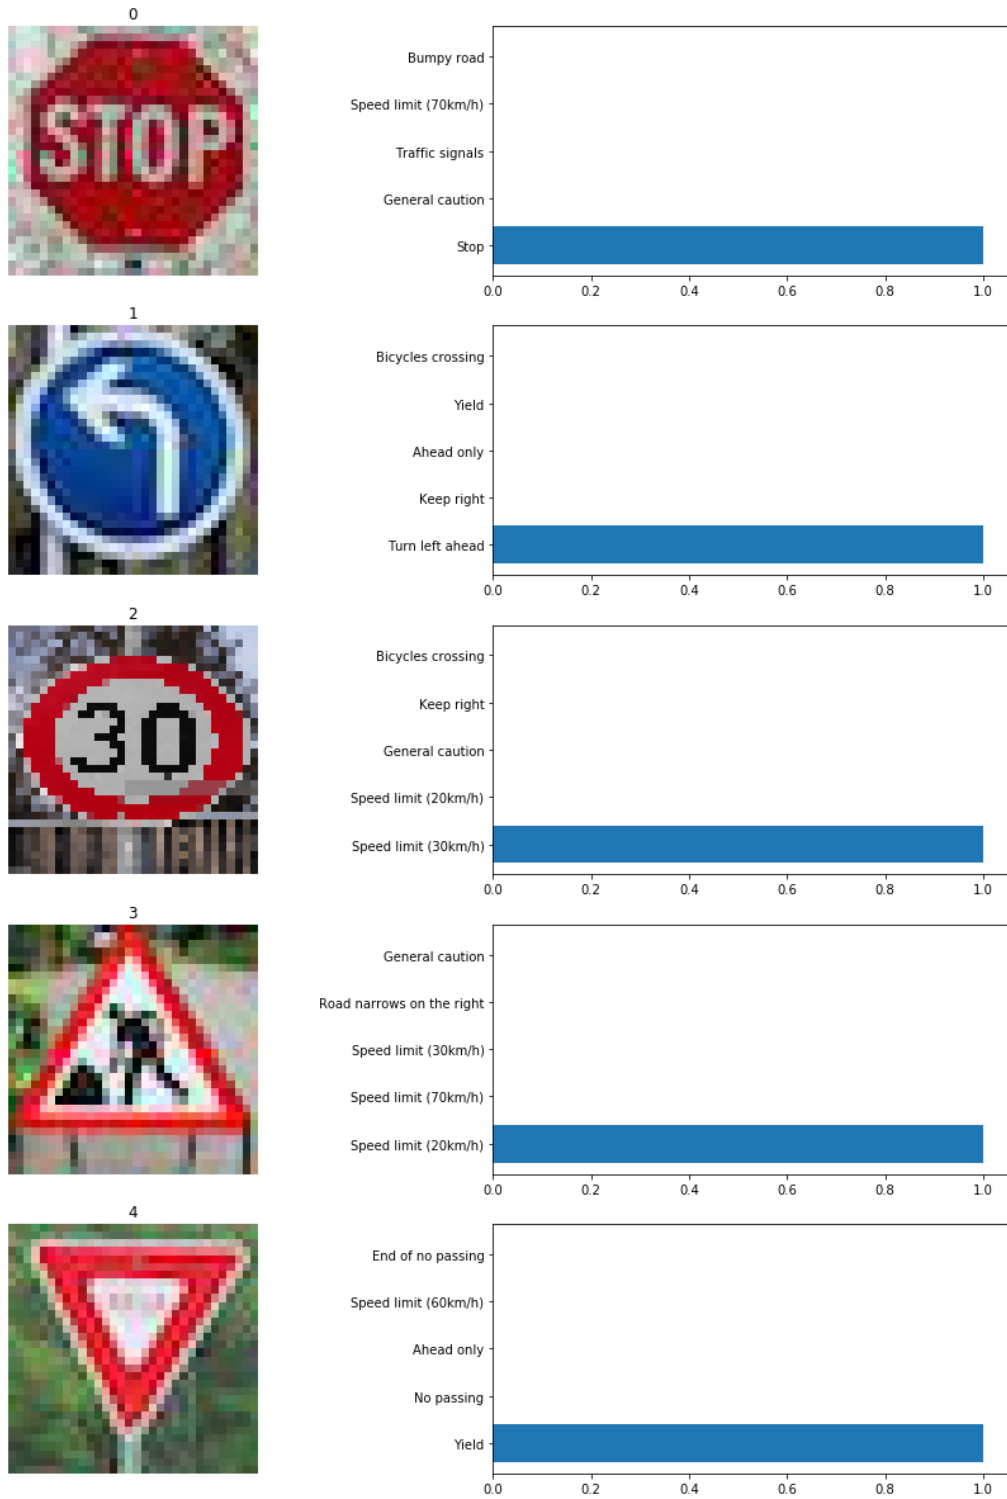


Figure 8 Top 5 Softmax probabilities of prediction of test images

From Figure 8, it is seen that the model is quite confident about its predictions. Fourth image on Roadwork was predicted Speed limit 20 km/h which is worrying.